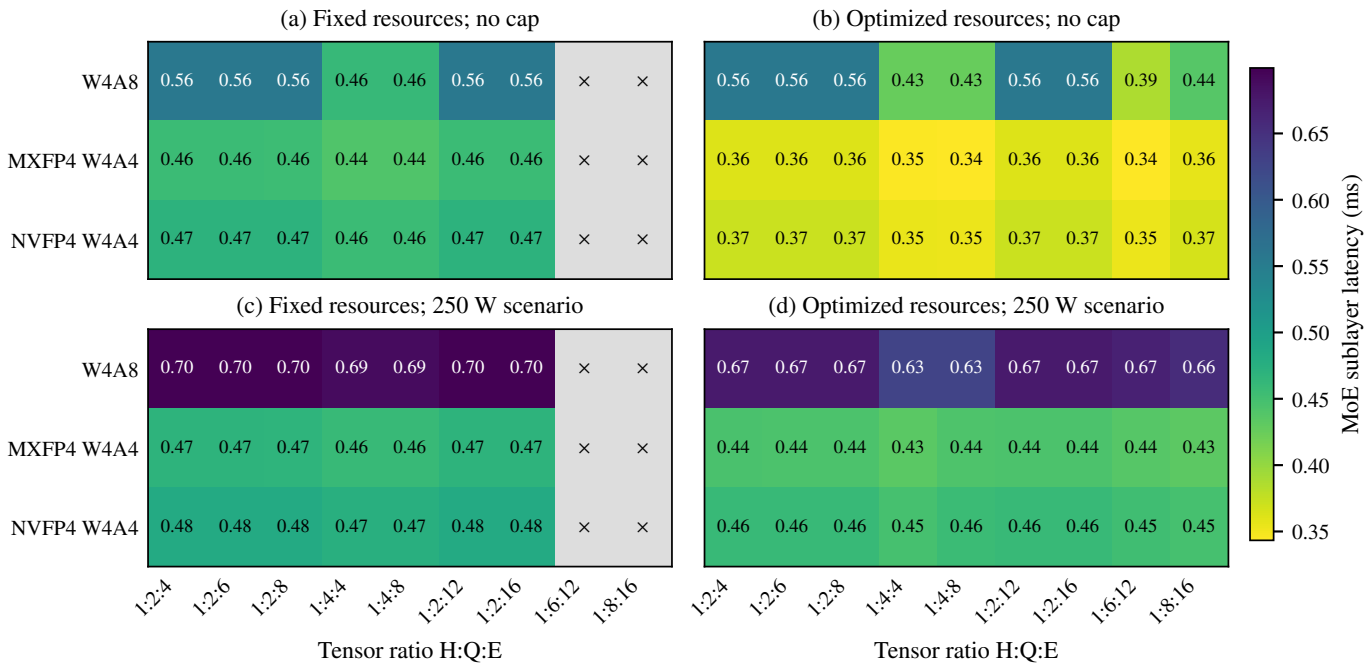
(a) Fixed resources; no cap
(b) Optimized resources; no cap
(c) Fixed resources; 250 W scenario
(d) Optimized resources; 250 W scenario
W4A8
MXFP4 W4A4
NVFP4 W4A4
0.56 0.56 0.56 0.46 0.46 0.56 0.56 × ×
0.46 0.46 0.46 0.44 0.44 0.46 0.46 × ×
0.47 0.47 0.47 0.46 0.46 0.47 0.47 × ×
0.56 0.56 0.56 0.43 0.43 0.56 0.56 0.39 0.44
0.36 0.36 0.36 0.35 0.34 0.36 0.36 0.34 0.36
0.37 0.37 0.37 0.35 0.35 0.37 0.37 0.35 0.37
0.70 0.70 0.70 0.69 0.69 0.70 0.70 × ×
0.47 0.47 0.47 0.46 0.46 0.47 0.47 × ×
0.48 0.48 0.48 0.47 0.47 0.48 0.48 × ×
0.67 0.67 0.67 0.63 0.63 0.67 0.67 0.67 0.66
0.44 0.44 0.44 0.43 0.44 0.44 0.44 0.44 0.43
0.46 0.46 0.46 0.45 0.46 0.46 0.46 0.45 0.45
1:2:4 1:2:6 1:2:8 1:4:4 1:4:8 1:2:12 1:2:16 1:6:12 1:8:16
Tensor ratio H:Q:E
MoE sublayer latency (ms)
0.65 0.60 0.55 0.50 0.45 0.40 0.35