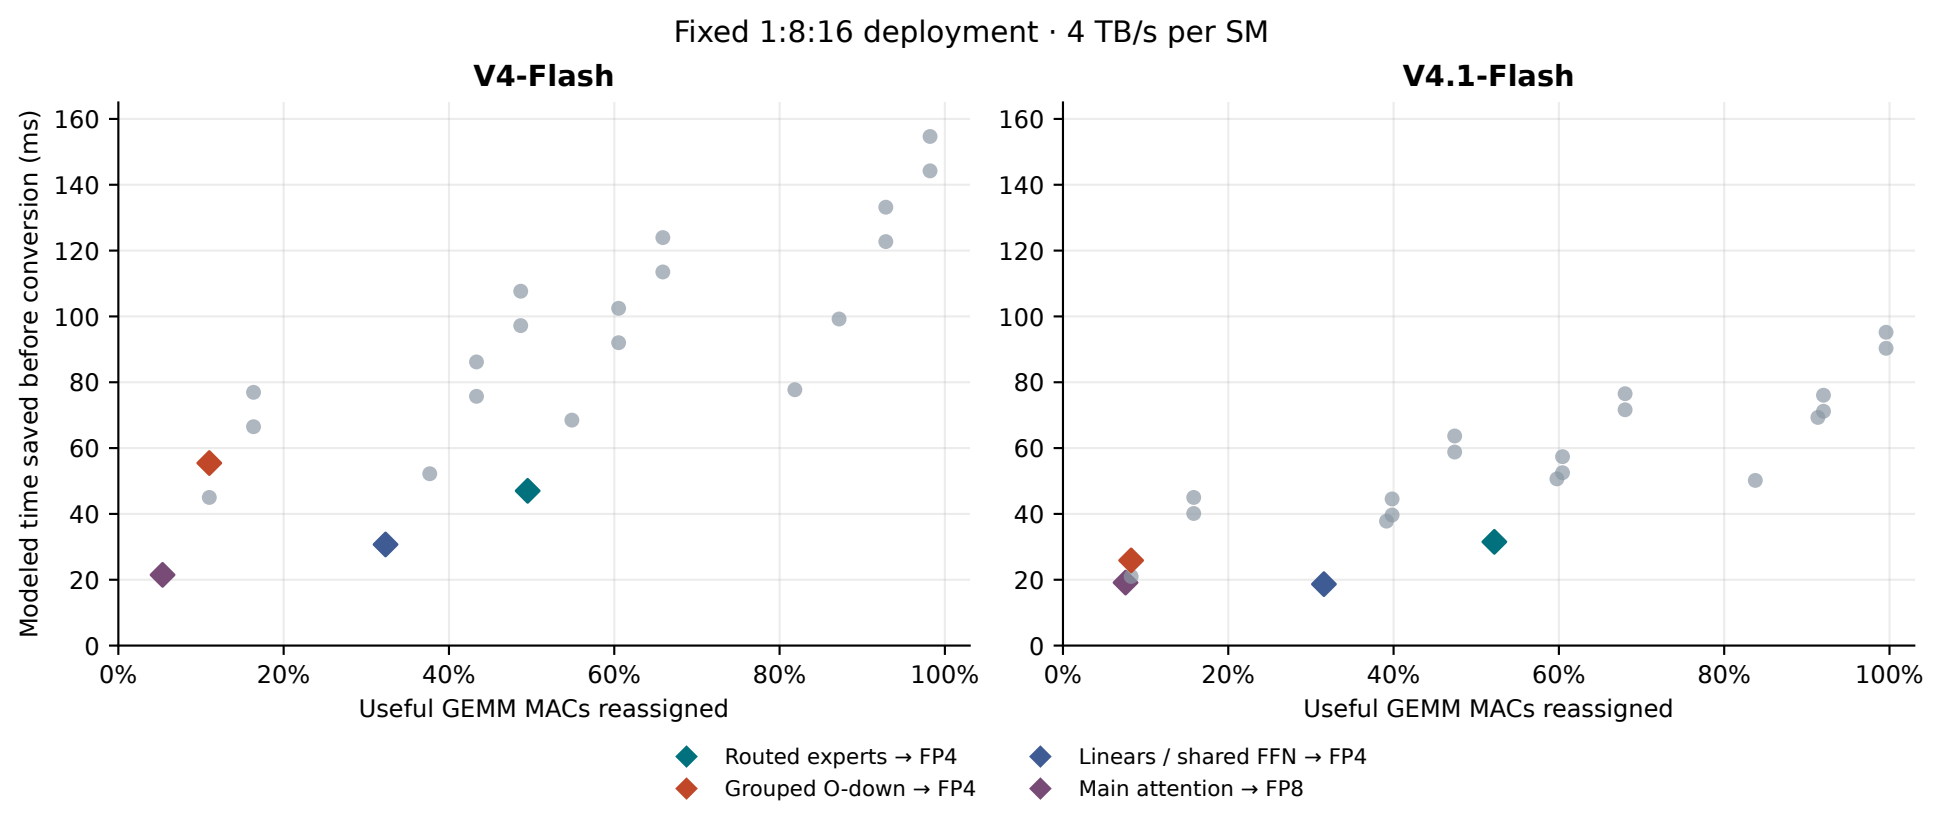
Fixed 1:8:16 deployment · 4 TB/s per SM
V4-Flash
V4.1-Flash
Modeled time saved before conversion (ms)
0
20
40
60
80
100
120
140
160
0%
20%
40%
60%
80%
100%
Useful GEMM MACs reassigned
Routed experts → FP4
Grouped O-down → FP4
Linears / shared FFN → FP4
Main attention → FP8